

# Performance Analysis Report: Go Load Balancer

Sikander

June 18, 2025

## Abstract

This report analyzes the performance of a Go-based load balancer designed to distribute requests across multiple FastAPI backends. An isolated performance test, bypassing backends, achieved 29,062 requests per second (req/s) with an average latency of 15.61ms. However, when proxying to backends, throughput dropped significantly to 500 req/s with one backend and 360 req/s with three. This report identifies bottlenecks, resource limitations, and logging overhead as primary causes of the performance degradation, providing detailed recommendations to optimize the system.

## 1 Introduction

The Go-based load balancer is implemented to distribute HTTP requests across multiple FastAPI backends using a round-robin algorithm, as specified in the configuration file (config.json). Initial tests showed a single FastAPI server handling 900 req/s. Introducing the load balancer reduced throughput to 500 req/s with one backend and 360 req/s with three backends, contrary to expectations of improved performance. An isolated test, bypassing backends, was conducted to measure the load balancer's standalone capacity, yielding 29,062 req/s. This report examines these results, identifies bottlenecks, and proposes solutions.

## 2 Test Configuration and Results

### 2.1 Isolated Test Setup

To isolate the load balancer's performance, the Balance function in balancer/balancer.go was modified to return a direct 200 OK response, bypassing backend proxying:

```
1 func (lb *LoadBalancer) Balance(w http.ResponseWriter, r *http.Request) {  
2     w.WriteHeader(http.StatusOK)  
3 }
```

The test was conducted using wrk with the following parameters:

- Threads: 4
- Connections: 500
- Duration: 10 seconds
- Target: http://localhost:8080

### 2.2 Results

The wrk output is summarized in Table 1:

Key observations:

Table 1: Load Balancer Isolated Test Results

Metric	Value
Total Requests	293,274
Duration	10.09 seconds
Requests per Second	29,062.24
Average Latency	15.61ms
Latency StdDev	16.10ms
Maximum Latency	338.97ms
Data Transferred	20.98MB (2.08MB/s)
Socket Errors	Connect: 0, Read: 0, Write: 243, Timeout: 0

- Throughput: The load balancer handled 29,062 req/s, indicating strong standalone performance.
- Latency: Average latency was 15.61ms, but a high standard deviation (16.10ms) and maximum latency (338.97ms) suggest variability.
- Errors: 243 write errors indicate potential overload with 500 connections.

### 3 Performance Analysis

#### 3.1 Comparison with Backend Tests

Previous tests showed:

- Single FastAPI Server: 900 req/s
- Load Balancer + 1 Backend: 500 req/s
- Load Balancer + 3 Backends: 360 req/s

The isolated test's 29,062 req/s confirms the load balancer's capacity is not the limiting factor. The significant drop when proxying suggests bottlenecks in the proxying process or backend interactions.

#### 3.2 Identified Bottlenecks

##### 3.2.1 Synchronous Logging

The load balancer code includes `log.Printf` statements in the request path, such as in the `Director` function of the reverse proxy:

```
1 log.Printf("Proxying request: %s %s Headers: %+v", req.Method,
    req.URL.String(), req.Header)
```

Synchronous logging is I/O-bound and can significantly slow down request processing under high load, contributing to the drop from 900 req/s to 500 req/s.

##### 3.2.2 Proxying Overhead

The `httputil.ReverseProxy` implementation incurs overhead from:

- Establishing connections to backends.
- Handling request and response transformations.

- Health checks every 100 seconds, which may not detect backend issues promptly.

This overhead likely exacerbates the performance drop when adding more backends.

### 3.2.3 Resource Contention

Running the load balancer and three backends on a single machine leads to competition for CPU, memory, and network resources. The isolated test's high throughput (29,062 req/s) pushed system resources, as evidenced by 243 write errors and high maximum latency (338.97ms).

## 3.3 Resource Limitations

The test was conducted on a single machine (DESKTOP-2N24STS), with unspecified hardware. Key limitations include:

- CPU: With 4 threads and 500 connections, each thread handled 125 connections. High CPU usage could cause latency spikes.
- Memory: Frequent logging and connection management may strain memory.
- Network: Although on localhost, high connection rates could saturate internal network buffers, contributing to write errors.

## 4 Recommendations

### 4.1 Optimize Logging

Remove or minimize synchronous logging in the request path. Modify the Director function to log only errors or use an asynchronous logging library like zap. Example:

```

1 proxy.Director = func(req *http.Request) {
2     req.URL.Scheme = u.Scheme
3     req.URL.Host = u.Host
4     req.Host = u.Host
5     req.URL.Path = path.Clean(req.URL.Path)
6     if req.URL.Path == "" || req.URL.Path == "." {
7         req.URL.Path = "/"
8     }
9     // Log only errors asynchronously
10 }
```

### 4.2 Enhance Proxy Efficiency

- Enable Keep-Alive: Ensure the http.Transport reuses connections:

```

1     proxy.Transport = &http.Transport{
2         DialContext: (&net.Dialer{Timeout: 5 *
3             time.Second}).DialContext,
4         ResponseHeaderTimeout: 30 * time.Second,
5         MaxIdleConns: 100,
6         IdleConnTimeout: 90 * time.Second,
7     }
```

- Faster Health Checks: Reduce the health check interval (e.g., to 10 seconds) in health/healthcheck.go.
- Load Balancing Algorithm: Consider switching to least\_conn if backends have varying performance.

### 4.3 Mitigate Resource Contention

- Monitor Resources: Use `htop` or `top` during tests to identify CPU or memory bottlenecks.
- Increase Threads: Adjust `wrk` to use more threads (e.g., 8) or tune `GOMAXPROCS` in Go.
- Distribute Components: Run backends on separate machines to reduce contention.

### 4.4 Further Testing

- Backend Isolation: Test each FastAPI backend individually to confirm their capacity (e.g., 900 req/s).
- Reduced Concurrency: Rerun the isolated test with fewer connections (e.g., 200) to check if write errors disappear.
- Alternative Load Balancers: Compare performance with Nginx or HAProxy to validate the Go implementation.

## 5 Conclusion

The isolated test demonstrated that the Go load balancer can handle 29,062 req/s with an average latency of 15.61ms, confirming its high standalone capacity. However, synchronous logging, proxying overhead, and resource contention caused significant performance degradation when proxying to backends (500 req/s with one backend, 360 req/s with three). By optimizing logging, enhancing proxy efficiency, and mitigating resource contention, the system can likely achieve or exceed the original 900 req/s. Implementing these recommendations and conducting further tests will ensure a robust and scalable load balancing solution.